

Legal Issues to Consider

Ethics

Morals - Ethics

Values

Principles

The Golden Rule


Universality

Transitivity

We, the members of the IEEE, in recognition of the importance of our technologies in affecting the quality of life throughout the world, and in accepting a personal obligation to our profession, its members and the communities we serve, do hereby commit ourselves to the highest ethical and professional conduct and agree:

1. to accept responsibility in making engineering decisions consistent with the safety, health and welfare of the public, and to disclose promptly factors that might endanger the public or the environment;
2. to avoid real or perceived conflicts of interest whenever possible, and to disclose them to affected parties when they do exist;
3. to be honest and realistic in stating claims or estimates based on available data;
4. to reject bribery in all its forms;
5. to improve the understanding of technology, its appropriate application, and potential consequences;
6. to maintain and improve our technical competence and to undertake technological tasks for others only if qualified by training or experience, or after full disclosure of pertinent limitations;
7. to seek, accept, and offer honest criticism of technical work, to acknowledge and correct errors, and to credit properly the contributions of others;
8. to treat fairly all persons regardless of such factors as race, religion, gender, disability, age, or national origin;
9. to avoid injuring others, their property, reputation, or employment by false or malicious action;
10. to assist colleagues and co-workers in their professional development and to support them in following this code of ethics.

DILBERT® by Scott Adams



© UFS, Inc.

Origins:

- Large investments in pipeline
- Tensions between engineering and marketing

Incorporate into the design:

- No single solution to the problem
- Failure mode analysis
- Apply safety standards
- Redundancy for reliability
- Inhibit misuse
- Tight cost/schedule estimates
- Design reviews

Intellectual Property & Legal Liability

Patents

Trade Secrets

Copyrights

Patents



- Citation of prior art
- Operation/reduce to practice
- Claims

Trade Secrets

Copyrights


(10) Patent No.: US 6,779,013 B2
(45) Date of Patent: Aug. 17, 2004

- | | | | | |
|-----------|----|-----------|----------------|----------|
| 5,764,089 | A | 6/1998 | Patovai et al. | 327/2001 |
| 5,838,730 | A | 4/1999 | Klass | 327/210 |
| 5,000,799 | A | 5/1990 | Tam | 327/201 |
| 5,993,051 | A | * 11/1999 | Jiang et al. | 708/505 |
| 6,205,462 | BI | * 3/2001 | Wyland et al. | 708/503 |
| 6,360,189 | A | * 3/2002 | Hinds et al. | 708/501 |
| 6,489,872 | BI | * 11/2002 | Chongtue | 708/505 |

Goto, G., et al., "A 54x54-b Regularly Structured Tree Multiplier", *IEEE Journal of Solid-State Circuits*, vol. 27, 1229-1236, (Sep. 1992).

A multiply-accumulate circuit includes a multiplier that generates a product with a binary exponent and a carry-save format. The product is converted to a normalized format having a three bit exponent and a fifty-bit mantissa. The multiplier generates a carry-save format for accumulation, and the multiplier also multiplies the converted products. Because the products being summed are in a carry-save format, post-normalization is avoided. The multiplier avoids a feedback loop. The adder operates on representations having exponents with a weight of thirty-two, and exponent and adder exponent paths are limited in size. The multiplier avoids the adder mantissa path. The multiplier generates a carry of thirty-two bits is provided by a carry-out.

20 Claims, 10 Drawings



(10) Patent No.: US 6,772,615 B2
(45) Date of Patent: Aug. 17, 2004

- | | | | | |
|-----------|----|-----------|---------------------|----------|
| 5,764,089 | A | 6/1998 | Pantovi et al. | 327/2000 |
| 5,898,330 | A | 4/1999 | Klass | 327/2100 |
| 5,900,759 | A | 5/1999 | Tam | 327/2011 |
| 5,993,051 | A | * 11/1999 | Jiang et al. | 708/5010 |
| 6,205,462 | B1 | * 3/2001 | Wyland et al. | 708/5030 |
| 6,360,189 | B1 | * 3/2002 | Hinds et al. | 708/5010 |
| 6,480,872 | B1 | * 11/2002 | Choquette | 708/5010 |

- (List continued on next page.)

Primary Examiner—Tan V. Mai
(74) **Attorney, Agent, or Firm**—Schwegman, Lundberg,
Wosessner & Kluth, P.A.

(57) ABSTRACT

A multiply-accumulate circuit includes a compressor tree to generate a product with a binary exponent and a mantissa in carry-save format. The product is converted into a number having a three bit exponent and a fifty-seven bit mantissa in carry-save format for accumulation. An adder circuit accu-

1

FLOATING POINT OVERFLOW AND SIGN DETECTION

FIELD

The present invention relates generally to floating point operations, and more specifically to floating point multiply accumulators.

BACKGROUND

Fast floating point mathematical operations have become an important feature in modern electronics. Floating point units are useful in applications such as three-dimensional graphics computations and digital signal processing (DSP). Examples of three-dimensional graphics computation include geometry transformations and perspective transformations. These transformations are performed when the motion of objects is determined by calculating physical equations in response to interactive events instead of displaying precomputed data.

Many DSP operations, such as finite impulse response (FIR) filters, compute $2A(b)$, where a is $n-1$, and a and b are both single precision floating point numbers. This type of computation typically employs floating point multiply accumulators (FMAC) units which perform many multiplication operations and add the resulting products to give the final result. In these types of operations, fast FMAC units typically execute multiplications and additions in parallel without pipeline bubbles. One example FMAC unit is described in: Nobuhiro et al., "2.44-GLOPS 300-MHz Floating-Point Vector Processing Unit for High-Performance 3-D Graphics Computing," IEEE Journal of Solid State Circuits, Vol. 35, No. 7, July 2000.

The Institute of Electrical and Electronic Engineers (IEEE) has published an industry standard for floating point operations in the ANSI/IEEE Std 754-1985, IEEE Standard for Binary Floating-Point Arithmetic, IEEE, New York, 1985, hereinafter referred to as the "IEEE standard." A typical implementation for a floating point FMAC compliant with the IEEE standard is shown in FIG. 1. FMAC 100 implements a single precision floating point multiply and accumulate instruction "D=AsBjC," as an indivisible operation. As can be seen from FIG. 1, fast floating point multipliers and fast floating point adders are both important ingredients to make a fast FMAC.

Multiplicands A and B are received by multiplier 110, and the product is normalized in post-normalization block 120. Multiplicands A and B are typically in an IEEE standard floating point format, and post-normalization block 120 typically operates on (normalizes) the output of multiplier 110 to make the product conform to the same format. For example, when multiplicands A and B are IEEE standard single precision floating point numbers, post-normalization block 120 operates on the output from multiplier 110 so that adder 130 receives the product as an IEEE standard single precision floating point number.

Adder 130 adds the normalized product from post-normalization block 120 with the output from multiplier 140. Multiplexer 140 can choose between the number C and the previous sum on node 152. When the previous sum is used, FMAC 100 is performing a multiply-accumulate function. The output of adder 130 is normalized in post-normalization block 150 so that the sum on node 152 is in the standard format discussed above.

Adder 130 and post-normalization block 150 can be "non-pipelined," which means that an accumulation can be

2

performed in a single clock cycle. When non-pipelined, adder 130 and post-normalization block typically include sufficient logic to limit the frequency at which FMAC 100 can operate, in part because floating point adders typically include circuits for alignment, mantissa addition, rounding, and other complex operations. To increase the frequency of operation, adder 130 and post-normalization block 150 can be "pipelined," which means registers can be included in the data path to store intermediate results. One disadvantage of pipelining is the introduction of pipeline stalls or bubbles, which decrease the effective data rate through FMAC 100.

For the reasons stated above, and for other reasons stated below which will become apparent to those skilled in the art upon reading and understanding the present specification, there is a need in the art for fast floating point multiply and accumulate circuits.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 shows a prior art floating point multiply-accumulate circuit;

FIG. 2 shows an integrated circuit with a floating point multiply-accumulate circuit;

FIG. 3 shows the exponent and mantissa paths of a floating point multiply-accumulate circuit;

FIG. 4 shows a mantissa multiplier circuit;

FIG. 5 shows a floating point conversion unit;

FIG. 6 shows a carry-save negation circuit;

FIG. 7 shows a base 32 floating point number representation;

FIG. 8 shows an exponent path of a floating point adder;

FIG. 9 shows a mantissa path of a floating point adder;

FIG. 10 shows an overflow detection circuit; and

FIG. 12 shows a sign detection circuit.

DESCRIPTION OF EMBODIMENTS

In the following detailed description of the embodiments, reference is made to the accompanying drawings which show, by way of illustration, specific embodiments in which the invention may be practiced. In the drawings, like numerals describe substantially similar components throughout the several views. These embodiments are described in sufficient detail to enable those skilled in the art to practice the invention. Other embodiments may be utilized and structural, logical, and electrical changes may be made without departing from the scope of the present invention. Moreover, it is to be understood that the various embodiments of the invention, although different, are not necessarily mutually exclusive. For example, a particular feature, structure, or characteristic described in one embodiment may be included within other embodiments. The following detailed description is, therefore, not to be taken in a limiting sense, and the scope of the present invention is defined only by the appended claims, along with the full scope of equivalents to which such claims are entitled.

Floating Point Multiplier Accumulator

FIG. 2 shows an integrated circuit with a floating point multiply-accumulate circuit. Integrated circuit 200 includes floating point multiplier 210, floating point conversion unit 220, floating point adder 230, and post-normalization circuit 240. Each of the elements shown in FIG. 2 is explained in further detail with reference to figures that follow. In this section, a brief overview of the FIG. 2 elements and their

11

-continued-

S1	C1	MC	Sign
1	0	0	0
1	0	1	1
1	1	1	1

Magnitude comparator 325 operates in parallel with adder mantissa path 324, so MC is available for sign detection circuit 1104 at substantially the same time as Msum. In this manner, the operation of sign detection circuit 1104 does not appreciably increase the delay within the feedback loop.

CONCLUSION

The method and apparatus of the present invention provide a fast multiply-accumulate operation that can be made compliant with any floating point format. Furthermore, the method and apparatus of the present invention can provide precision comparable to the precision available using prior art double precision arithmetic units, in part because the mantissa fields are expanded. In some embodiments, IEEE standard single precision operands are multiplied and the products are summed. The multiplier includes a compressor tree to generate a product with a binary exponent and a mantissa in carry-save format. The product is converted into a number having a three bit exponent and a fifty-seven bit mantissa in carry-save format for accumulation. An accumulator circuit accumulates the converted products in carry-save format. Because the products being summed are in carry-save format, post-normalization is avoided within the adder feedback loop. In addition, because the adder operates on floating point number representations having exponents with a least significant bit weight of thirty-two, exponent comparisons within the adder exponent path are fast, and variable shifters can be avoided in the adder mantissa path. When the adder is not pipelined, a fast single cycle accumulation is realized with the method and apparatus of the present invention.

It is to be understood that the above description is intended to be illustrative, and not restrictive. Many other embodiments will be apparent to those of skill in the art upon reading and understanding the above description. The scope of the invention should, therefore, be determined with reference to the appended claims, along with the full scope of equivalents to which such claims are entitled.

What is claimed is:

1. A floating point multiply-accumulate circuit comprising:

an exponent path including:
an exponent summer to sum two input exponents having a first weight to produce a product exponent;
an exponent conversion unit coupled to the output of the exponent summer, to convert the product exponent to a second weight; and
an exponent accumulation stage to accumulate the converted product exponent and to choose a larger exponent from the converted product exponent and an accumulated exponent; and

a mantissa path including:
a mantissa multiplier to multiply two input mantissas and produce a product mantissa in carry-save format that includes a sum field and a carry field;
a mantissa shifter to shift the sum field and the carry field responsive to the exponent conversion unit in the exponent path; and

12

a mantissa accumulator to accumulate shifted product mantissas in carry-save format, the mantissa accumulator including an overflow detection circuit responsive to two most significant bits of a sum output from the mantissa accumulator.

2. The floating point multiply-accumulate circuit of claim 1, wherein the overflow detection circuit comprises exclusive-or gate.

3. The floating point multiply-accumulate circuit of claim 1, wherein the exponent conversion unit is configured to convert the least significant five bits of the product exponent.

4. The floating point multiply-accumulate circuit of claim 1, wherein the mantissa shifter is configured to shift the field and carry field of the product mantissa by a number of bit positions equal to a value of the least significant five bits of the product exponent.

5. The floating point multiply-accumulate circuit of claim 1, wherein the mantissa accumulator comprises four-to-two compressors.

6. The floating point multiply-accumulate circuit of claim 1, further comprising a post-normalization stage to produce a normalized floating point resultant.

7. The floating point multiply-accumulate circuit of claim 1, wherein the post-normalization stage includes a sign detection circuit.

8. The floating point multiply-accumulate circuit of claim 1, further comprising a magnitude comparator with the mantissa accumulator, wherein the circuit is responsive to the magnitude comparator.

9. The floating point multiply-accumulate circuit of claim 1, wherein the post-normalization stage is configured to convert the product exponent to a second weight.

10. The floating point multiply-accumulate circuit of claim 1, wherein the exponent conversion unit is configured to convert the product exponent to have a least weight equal to thirty-two.

11. A method performed within a programmable computer, comprising:

multiplying two floating point mantissas;
two floating point exponents to form a product exponent;
converting the product to have a different bit weight exponent field;

accumulating the converted product in a carry-save format;

detecting overflow as a function of two most significant bits of a sum field of an accumulated product;

post-normalizing the accumulated product;

12. The method of claim 11 wherein said product comprises adding a first plurality of said product, the method further comprising post-normalizing until the last product is a product of two.

13. The method of claim 11 wherein said product comprises:

shifting a mantissa of the product by an amount equal to the value of the least significant five bits of the product; and

zeroing the least significant five bits of the product.

14. The method of claim 11 wherein said product comprises:

comparing an exponent of a first converted product of a second converted product; and
conditionally shifting right by a fixed amount of the converted product having a smaller selecting the larger exponent as a resultant product.

What is claimed is:

1. A floating point multiply-accumulate circuit comprising:

an exponent path including:

an exponent summer to sum two input exponents having a first weight to produce a product exponent;
an exponent conversion unit coupled to the output of the exponent summer, to convert the product exponent to a second weight; and
an exponent accumulation stage to accumulate the converted product exponent and to choose a larger exponent from the converted product exponent and an accumulated exponent; and

a mantissa path including:

a mantissa multiplier to multiply two input mantissas and produce a product mantissa in carry-save format that includes a sum field and a carry field;
a mantissa shifter to shift the sum field and the carry field responsive to the exponent conversion unit in the exponent path; and

a mantissa accumulator to accumulate shifted product mantissas in carry-save format, the mantissa accumulator including an overflow detection circuit responsive to two most significant bits of a sum field output from the mantissa accumulator.

2. The floating point multiply-accumulate circuit of claim 1, wherein the overflow detection circuit comprises an exclusive-or gate.

3. The floating point multiply-accumulate circuit of claim 1, wherein the exponent conversion unit is configured to convert the least significant five bits of the product exponent.

4. The floating point multiply-accumulate circuit of claim 1, wherein the mantissa shifter is configured to shift the sum field and carry field of the product mantissa by a number of bit positions equal to a value of the least significant five bits of the product exponent.

5. The floating point multiply-accumulate circuit of claim 1, wherein the mantissa accumulator comprises four-to-two compressors.

6. The floating point multiply-accumulate circuit of claim 1, further comprising a post-normalization stage to produce a normalized floating point resultant.

7. The floating point multiply-accumulate circuit of claim 1, wherein the post-normalization stage includes a sign detection circuit.

8. The floating point multiply-accumulate circuit of claim 1, further comprising a magnitude comparator in parallel with the mantissa accumulator, wherein the sign detection circuit is responsive to the magnitude comparator.

9. The floating point multiply-accumulate circuit of claim 1, wherein the post-normalization stage is configured to be turned off until accumulation is complete.

10. The floating point multiply-accumulate circuit of claim 1, wherein the exponent conversion unit is configured to convert the product exponent to have a least significant bit weight equal to thirty-two.

Liability / Negligence

Product Liability
(Product Design)



Negligence vs. Strict liability

Negligence

Did not follow
reasonable standards

Strict Liability

Dangerous or defective
at the time of
manufacture